



# Amdahl's Law

Amdahl's Law Identifies performance gains from adding additional cores to an application that has both serial and parallel components

S is serial portion

N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times

If we have 2 cores ( $p = 2$ ), the speedup would be:

$$S_{2 \text{ cores}} = \frac{1}{0.25 + \frac{0.75}{2}} = \frac{1}{0.625} \approx 1.6$$





# Amdhal's Law

---

$$F_s = 0.20$$

$$F_p = 0.80$$

$$N = 5$$

$$\text{Speedup} = \frac{1}{(0.20 + \frac{0.80}{5})}$$

$$\text{Speedup} = \frac{1}{(0.20 + 0.16)}$$

$$\text{Speedup} = \frac{1}{0.36}$$

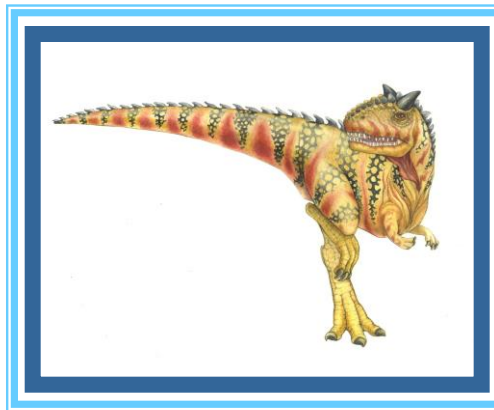
$$\text{Speedup} \approx 2.78$$

So, the potential speedup with 5 threads is approximately 2.78.



# Chapter 7: Main Memory

---





# Chapter 7: Memory Management

---

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table





# Objectives

---

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation





# Memory Management Basics

"**Memory** is the faculty of the brain by which **data or information** is **encoded, stored, and retrieved** when needed."

-- Wikipedia





# Memory Management Basics

"**Memory** is the faculty of the brain by which **data** or **information** is **encoded, stored, and retrieved** when needed."

-- Wikipedia





# Memory Management Basics

"Memory is the faculty of the brain by which data or information is encoded, stored, and retrieved when needed."

-- Wikipedia








# Memory Management Basics

"**Memory** is the faculty of the brain by which **data or information** is **encoded, stored, and retrieved** when needed."

-- Wikipedia





Frequency: 2 GHz

$$\text{Time: } \frac{1}{\text{Frequency}}$$
$$= \frac{1}{2 \times 10^9} \text{ Sec.}$$

1 Kilo <Unit>	= 1000 <Unit>	= $10^3$ <Unit>	
1 Mega <Unit>	= 1000 Kilo <Unit>	= $10^3 \times 10^3$ <Unit>	= $10^6$ <Unit>
1 Giga <Unit>	= 1000 Mega <Unit>	= $10^3 \times 10^3 \times 10^3$ <Unit>	= $10^9$ <Unit>





Frequency: 2 GHz

$$\begin{aligned}\text{Time: } & \frac{1}{\text{Frequency}} \\ &= \frac{1}{2 \times 10^9} \text{ Sec.} \\ &= \frac{1}{2} \times 10^{-9} \text{ Sec.} \\ &= \frac{1}{2} \text{ nsec.}\end{aligned}$$





# Memory Management Basics

HIERARCHY




RANKING






# Memory Management Basics





# Memory Management Basics



Nolan's Dark Knight  
Tribology: Release Date





# Memory Management Basics



Release Year



Tomato Meter: Heath Ledger



**Rotten Tomatoes**




**metacritic**








# Memory Management Basics





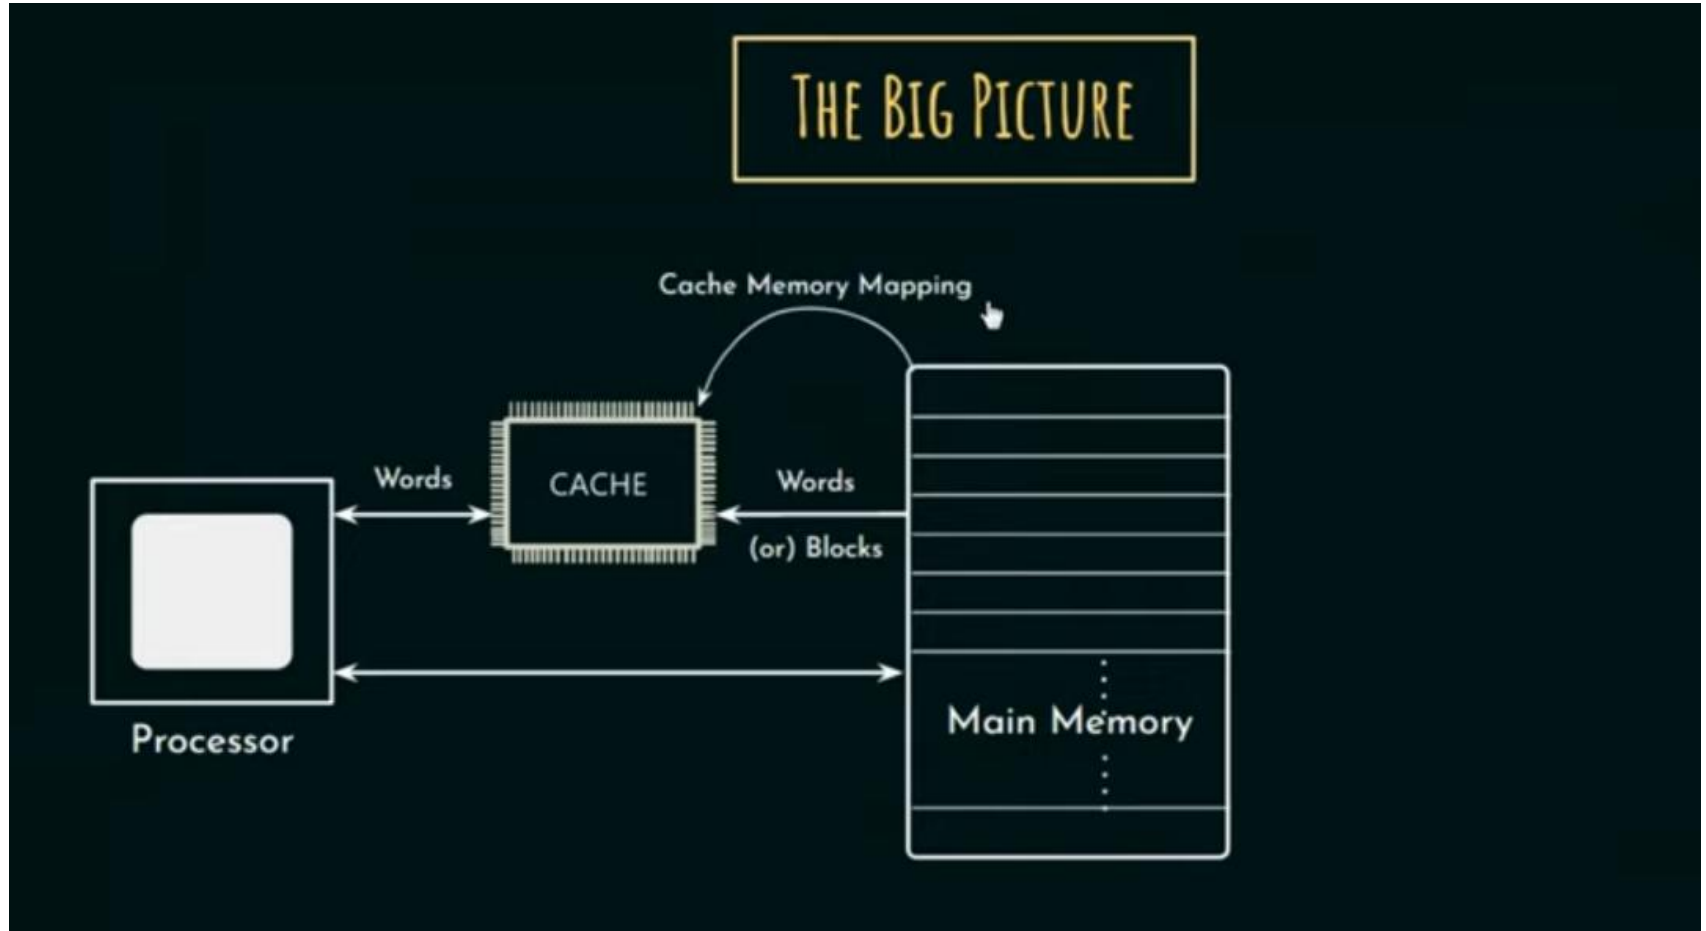


# Memory Management Basics



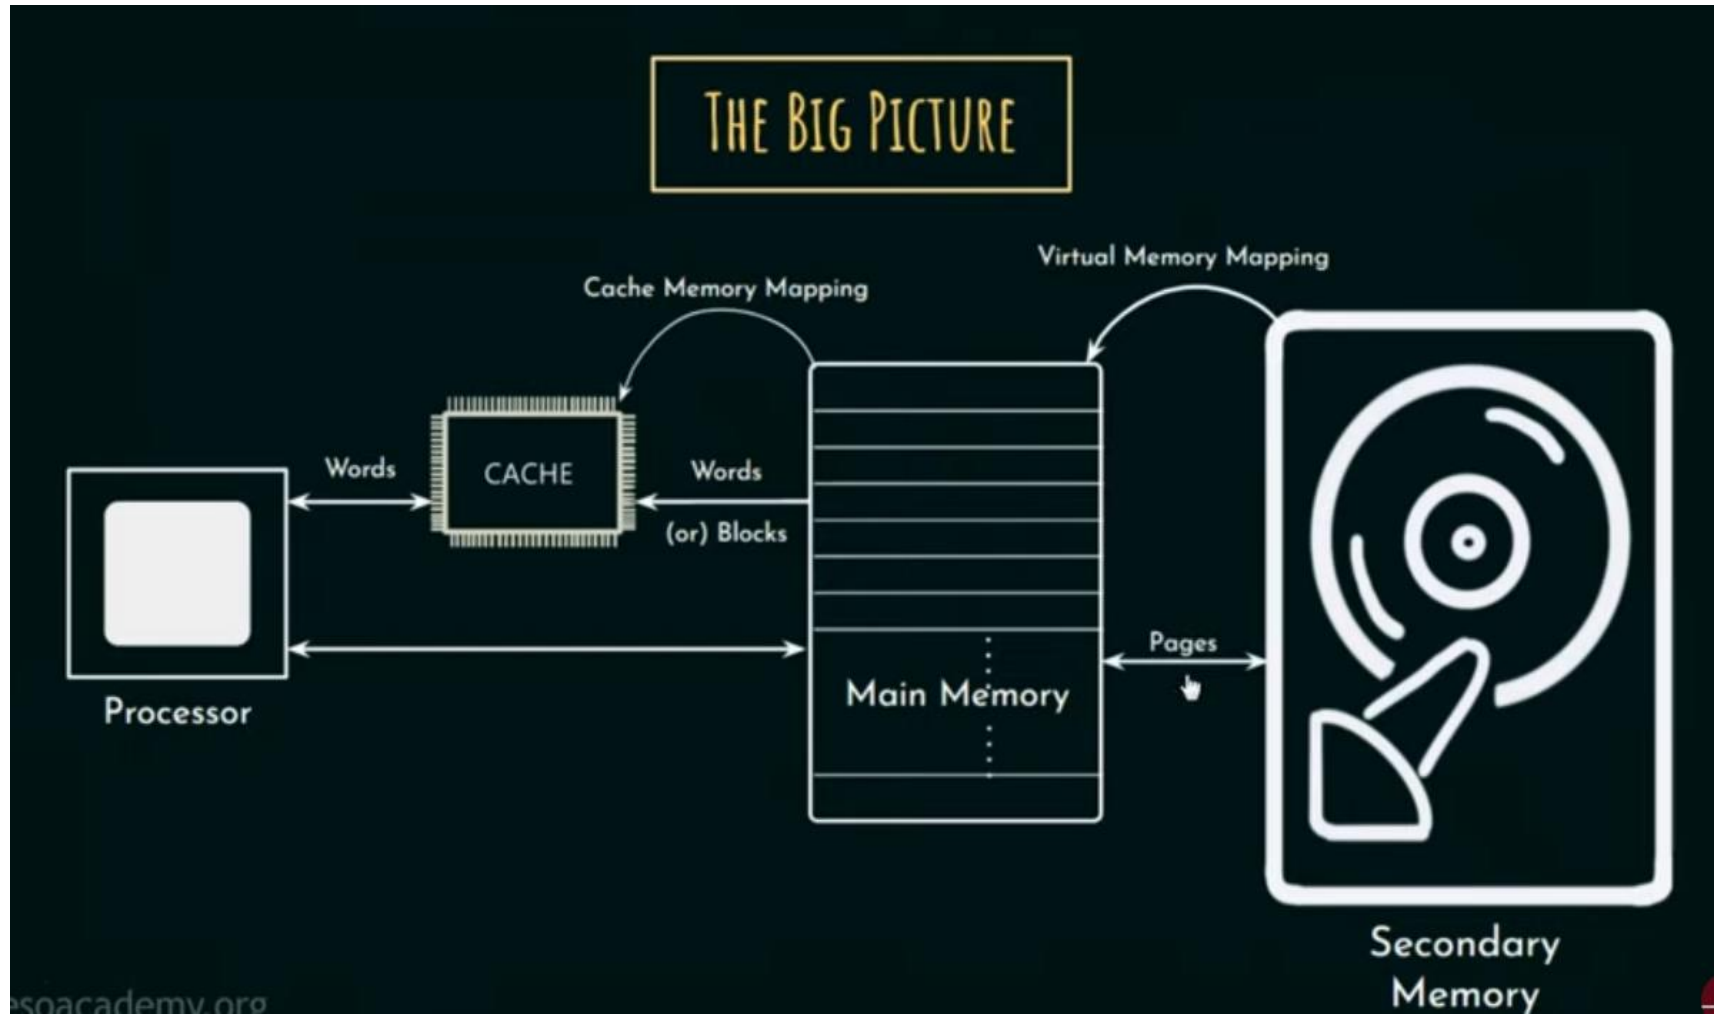


# Memory Management Basics





# Memory Management Basics





# Background

---


- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are only storage CPU can access directly
- Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- Register access in one CPU clock (or less)
- Main memory can take many cycles, causing a **stall**
- **Cache** sits between main memory and CPU registers
- Protection of memory required to ensure correct operation





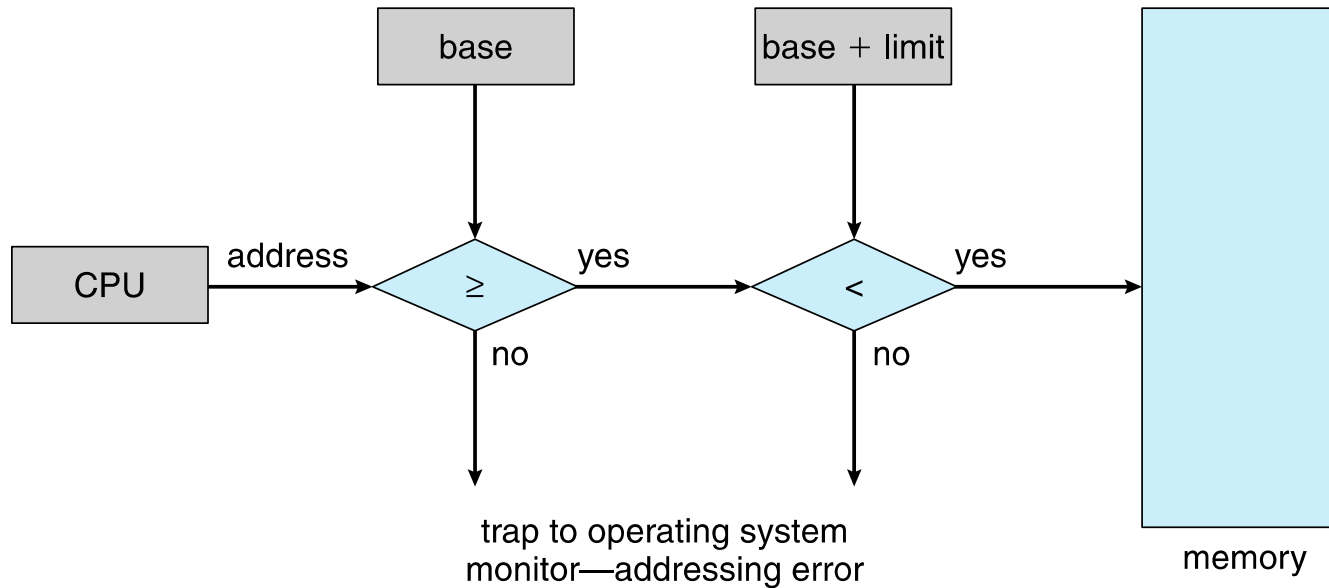
# Base and Limit Registers

- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user





# Hardware Address Protection





# Address Binding

- Programs on disk, ready to be brought into memory to execute form an **input queue**
  - Without support, must be loaded into address 0000
- Inconvenient to have first user process physical address always at 0000
  - How can it not be?
- Further, addresses represented in different ways at different stages of a program's life
  - Source code addresses usually symbolic
  - Compiled code addresses **bind** to relocatable addresses
    - ▶ i.e. “14 bytes from beginning of this module”
  - Linker or loader will bind relocatable addresses to absolute addresses
    - ▶ i.e. 74014
  - Each binding maps one address space to another





# Binding of Instructions and Data to Memory

---

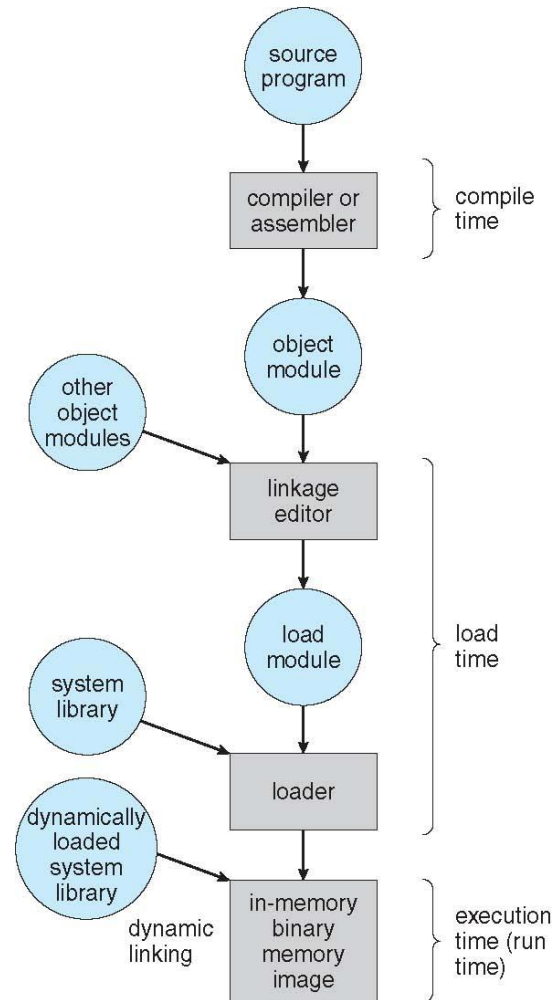
- Address binding of instructions and data to memory addresses can happen at three different stages
  - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
  - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
  - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
    - ▶ Need hardware support for address maps (e.g., base and limit registers)







# Multistep Processing of a User Program





# Logical vs. Physical Address Space

---

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
  - **Logical address** – generated by the CPU; also referred to as **virtual address**
  - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program





# Memory-Management Unit (MMU)

---


- Hardware device that at run time maps virtual to physical address
- Many methods possible, covered in the rest of this chapter
- To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
  - Base register now called **relocation register**
  - MS-DOS on Intel 80x86 used 4 relocation registers
- The user program deals with *logical* addresses; it never sees the *real* physical addresses
  - Execution-time binding occurs when reference is made to location in memory
  - Logical address bound to physical addresses





# Dynamic relocation using a relocation register

- Routine is not loaded until it is called
- Better memory-space utilization; unused routine is never loaded
- All routines kept on disk in relocatable load format
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
  - Implemented through program design
  - OS can help by providing libraries to implement dynamic loading





# Dynamic Linking

- **Static linking** – system libraries and program code combined by the loader into the binary program image
- Dynamic linking –linking postponed until execution time
- Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine, and executes the routine
- Operating system checks if routine is in processes' memory address
  - If not in address space, add to address space
- Dynamic linking is particularly useful for libraries
- System also known as **shared libraries**
- Consider applicability to patching system libraries
  - Versioning may be needed






# Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
  - Total physical memory space of processes can exceed physical memory
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk





# Schematic View of Swapping





# Swapping (Cont.)

---

- Does the swapped out process need to swap back in to same physical addresses?
- Depends on address binding method
  - Plus consider pending I/O to / from process memory space
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
  - Swapping normally disabled
  - Started if more than threshold amount of memory allocated
  - Disabled again once memory demand reduced below threshold



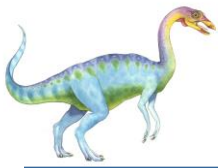




# Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
  - Swap out time of 2000 ms
  - Plus swap in of same sized process
  - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped – by knowing how much memory really being used
  - System calls to inform OS of memory use via `request_memory()` and `release_memory()`



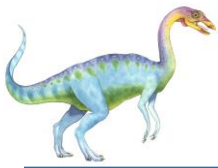


# Contiguous Allocation

---

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two **partitions**:
  - Resident operating system, usually held in low memory with interrupt vector
  - User processes then held in high memory
  - Each process contained in single contiguous section of memory





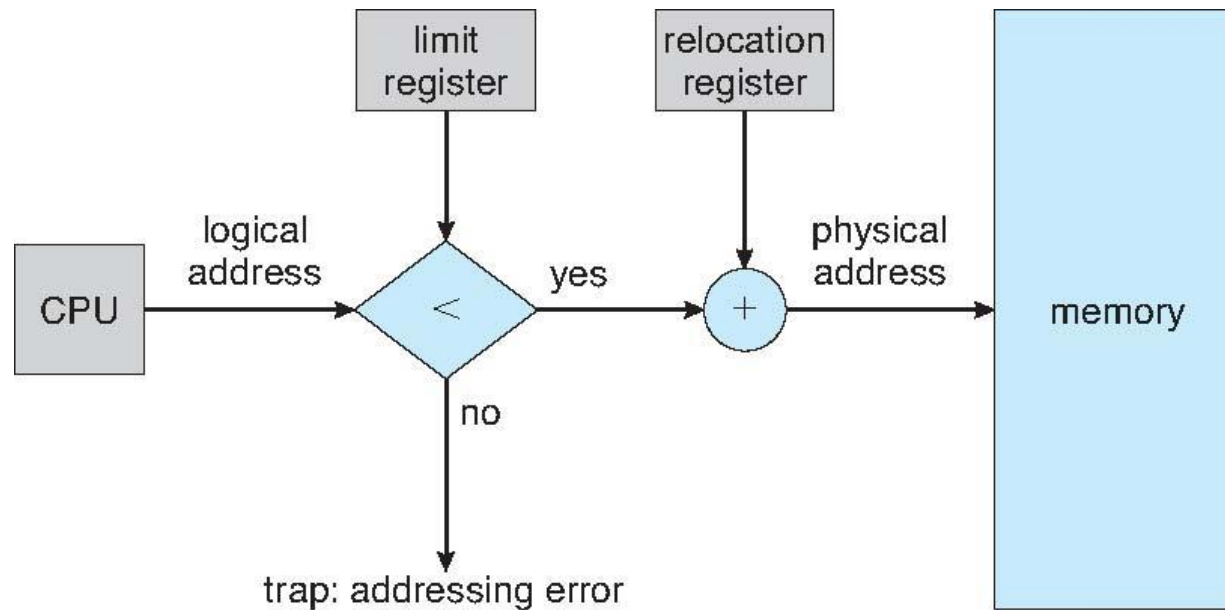
# Contiguous Allocation (Cont.)

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
  - Base register contains value of smallest physical address
  - Limit register contains range of logical addresses – each logical address must be less than the limit register
  - MMU maps logical address *dynamically*
  - Can then allow actions such as kernel code being **transient** and kernel changing size





# Hardware Support for Relocation and Limit Registers






# Multiple-partition allocation

## ■ Multiple-partition allocation

- Degree of multiprogramming limited by number of partitions
- **Variable-partition** sizes for efficiency (sized to a given process' needs)
- **Hole** – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Process exiting frees its partition, adjacent free partitions combined
- Operating system maintains information about:  
a) allocated partitions    b) free partitions (hole)





# Dynamic Storage-Allocation Problem

---

How to satisfy a request of size  $n$  from a list of free holes?

- **First-fit**: Allocate the **first** hole that is big enough
- **Best-fit**: Allocate the **smallest** hole that is big enough; must search entire list, unless ordered by size
  - Produces the smallest leftover hole
- **Worst-fit**: Allocate the **largest** hole; must also search entire list
  - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization





# Fragmentation

---

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- First fit analysis reveals that given  $N$  blocks allocated,  $0.5 N$  blocks lost to fragmentation
  - 1/3 may be unusable -> **50-percent rule**





# Fragmentation (Cont.)

---

- Reduce external fragmentation by **compaction**
  - Shuffle memory contents to place all free memory together in one large block







# Fragmentation

Fragmented memory before compaction



Memory after compaction






# Paging

- ❑ With paging, physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
  - ❑ Avoids external fragmentation
  - ❑ Avoids problem of varying sized memory chunks
- ❑ Divide physical memory into fixed-sized blocks called **frames**
  - ❑ Size is power of 2, between 512 bytes and 1 Gbytes
- ❑ Divide logical memory into blocks of same size called **pages**
- ❑ Keep track of all free frames
- ❑ To run a program of size **N** pages, need to find **N** free frames and load program
- ❑ Set up a **page table** to translate logical to physical addresses
- ❑ Backing store likewise split into pages
- ❑ Still have Internal fragmentation

64bits: Several TB of  
RAM: 40bits:TB,  
50:PB, 60:HB,


1KB: 2 Power10  
1MB: 2 Power20  
1GB: 2 Power30  
2GB: 2 Power31  
4GB: 2 Power32





# Address Translation Scheme

- Address generated by CPU is divided into:
  - **Page number** ( $p$ ) – used as an index into a **page table** which contains base address of each page in physical memory
  - **Page offset** ( $d$ ) – combined with base address to define the physical memory address that is sent to the memory unit




- For given logical address space  $2^m$  and page size  $2^n$






# Paging Hardware






# Paging Model of Logical and Physical Memory






# Paging Model of Logical and Physical Memory





# Paging Example

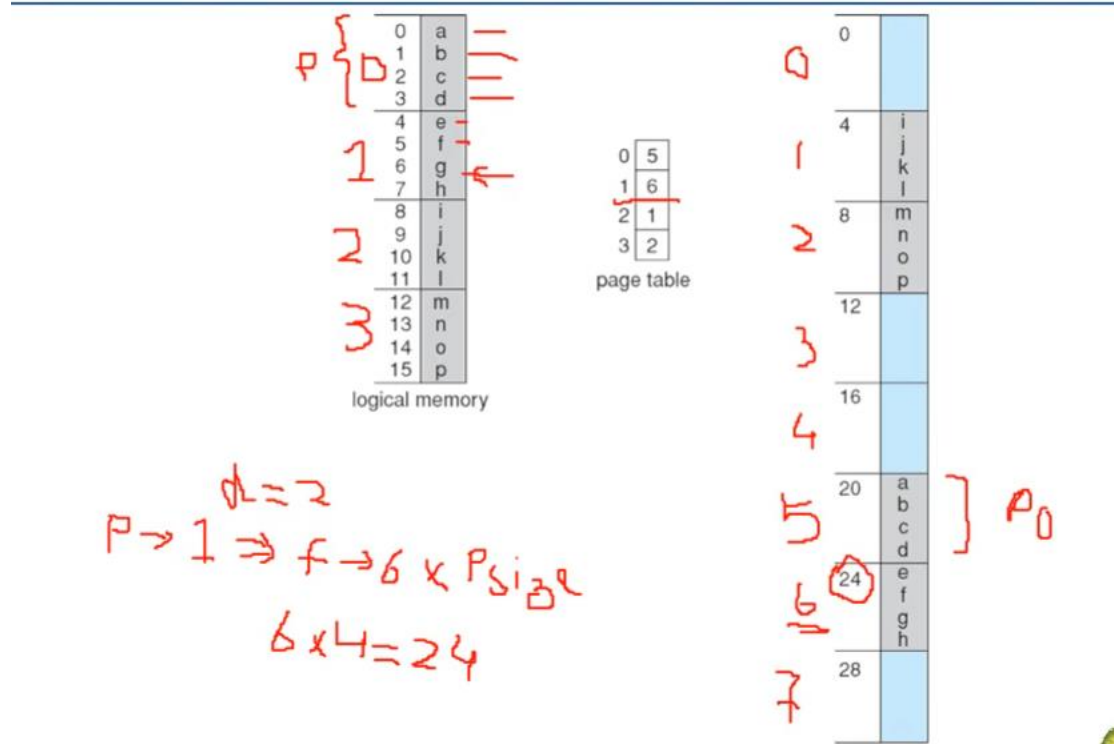


$n=2$  and  $m=4$  32-byte memory and 4-byte pages





# Paging Example







# Paging (Cont.)

## ■ Calculating internal fragmentation

- Page size = 2,048 bytes
- Process size = 72,766 bytes
- 35 pages + 1,086 bytes
- Internal fragmentation of  $2,048 - 1,086 = 962$  bytes
- Worst case fragmentation = 1 frame – 1 byte
- On average fragmentation =  $1 / 2$  frame size
- So small frame sizes desirable?
- But each page table entry takes memory to track
- Page sizes growing over time

Page Size: 2KB  
71,680: 35 Pages  
2049B: Eg2


- ▶ Solaris supports two page sizes – 8 KB and 4 MB

- Process view and physical memory now very different
- By implementation process can only access its own memory





# Free Frames



Before allocation

After allocation





# Segmentation


---

- Memory-management scheme that supports user view of memory
- A program is a collection of segments
  - A segment is a logical unit such as:
    - main program
    - procedure
    - function
    - method
    - object
    - local variables, global variables
    - common block
    - stack
    - symbol table
    - arrays





# User's View of a Program



$$\begin{array}{r} 2048B \times 2 \\ = 4096B \\ \hline 4096B \end{array}$$


$$\begin{array}{r} 2048B P_1 \\ \hline 2048B P_2 \\ 952B \\ \hline 1096B \end{array}$$

$$\begin{array}{r} P_2 = 3000B. \\ 2048B \\ \hline 952B \end{array}$$

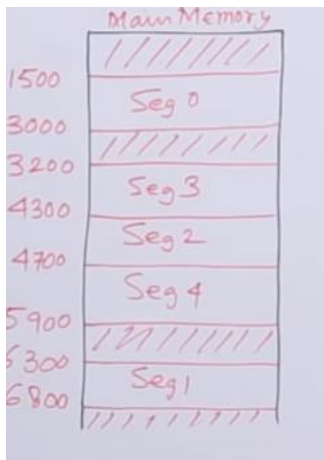




# User's View of a Program



	Limit	Base
0	1500	1500
1	500	6300
2	400	4300
3	1100	3200
4	1200	4700




logical address






# Logical View of Segmentation



user space



physical memory space





# Segmentation Architecture

- Logical address consists of a two tuple:  
    <segment-number, offset>,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
  - **base** – contains the starting physical address where the segments reside in memory
  - **limit** – specifies the length of the segment
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;  
    segment number **s** is legal if **s** < **STLR**





# Segmentation Architecture (Cont.)

---


- Protection
  - With each entry in segment table associate:
    - ▶ validation bit = 0  $\Rightarrow$  illegal segment
    - ▶ read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- A segmentation example is shown in the following diagram








# Segmentation Hardware



400  
 $d \rightarrow (0 \text{ to } 399)$





# Segmentation Hardware

Consider following segment table:

Segment No.	Base	Length
0	1219	700
1	2300	14
2	90	100
3	1327	580
4	1952	96

1. Physical address =  $1219 + 430$

2. " " =  $2300 + 11$

3. Trap addressing error

4. " " =  $1327 + 425$

5. " " =  $1952 + 95$

Which following logical address  
will produce Trap addressing error

1. 0, 430      2. 1, 11      3. 2, 100

4. 3, 425      5. 4, 95

Ans: A) 1      B) 2      C) 3      D) 5





# Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table: Extra Protection: Can generate Trap.
- In this scheme every data/instruction access requires two memory accesses
  - One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**
- **Cache Analogy**

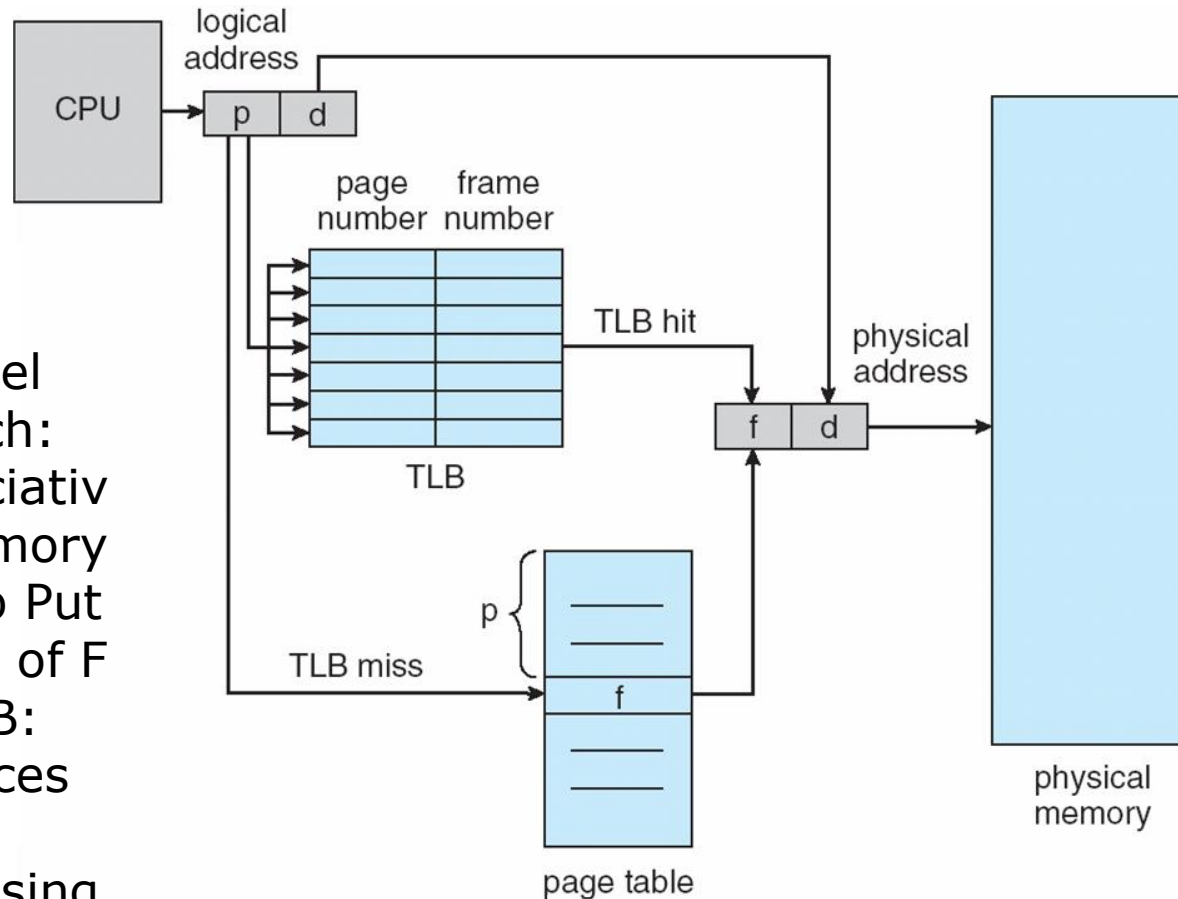
Location of Page table Pointed by





# Paging Hardware With TLB

1. Parallel Search: Associative Memory
2. Try to Put Value of F in TLB: Chances of accessing same page again.





# Associative Memory

- Associative memory – parallel search

Page #	Frame #

- Address translation (p, d)
  - If p is in associative register, get frame # out
  - Otherwise get frame # from page table in memory





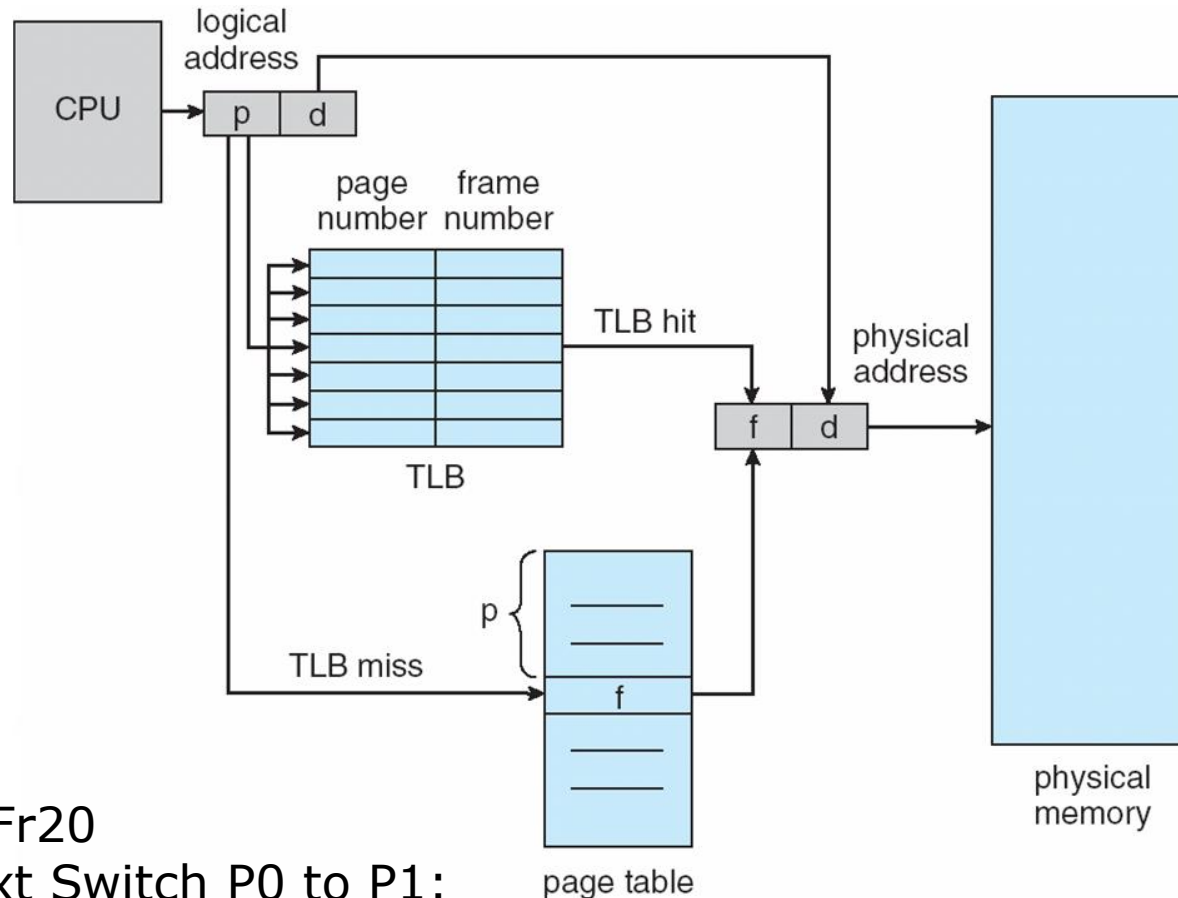
# Implementation of Page Table (Cont.)

- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
  - Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
  - Some entries can be **wired down** for permanent fast access





# Paging Hardware With TLB



- Pg15-Fr20
- Context Switch P0 to P1:
- Pg15-Fr20 in actual PT
- Stale Entry in TLB
- P0-Pg15-F2r0
- P0-Pg15-Fr20





# Effective Access Time

- Associative Lookup =  $\varepsilon$  time unit
    - Can be < 10% of memory access time
  - Hit ratio =  $\alpha$ 
    - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
  - Consider  $\alpha = 80\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - **Effective Access Time (EAT)**
    - 80% find entry in 1 Access TLB
    - 20% find entry 2Acces
- Consider  $\alpha = 80\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
- $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio : all imp pages get loaded in to TLB ->  $\alpha = 99\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$ : **Very close to actual MA: instead of paging overhead**

1. TLB Miss Happens
2. Example to Find Access time







# How Memory Protection: Paging

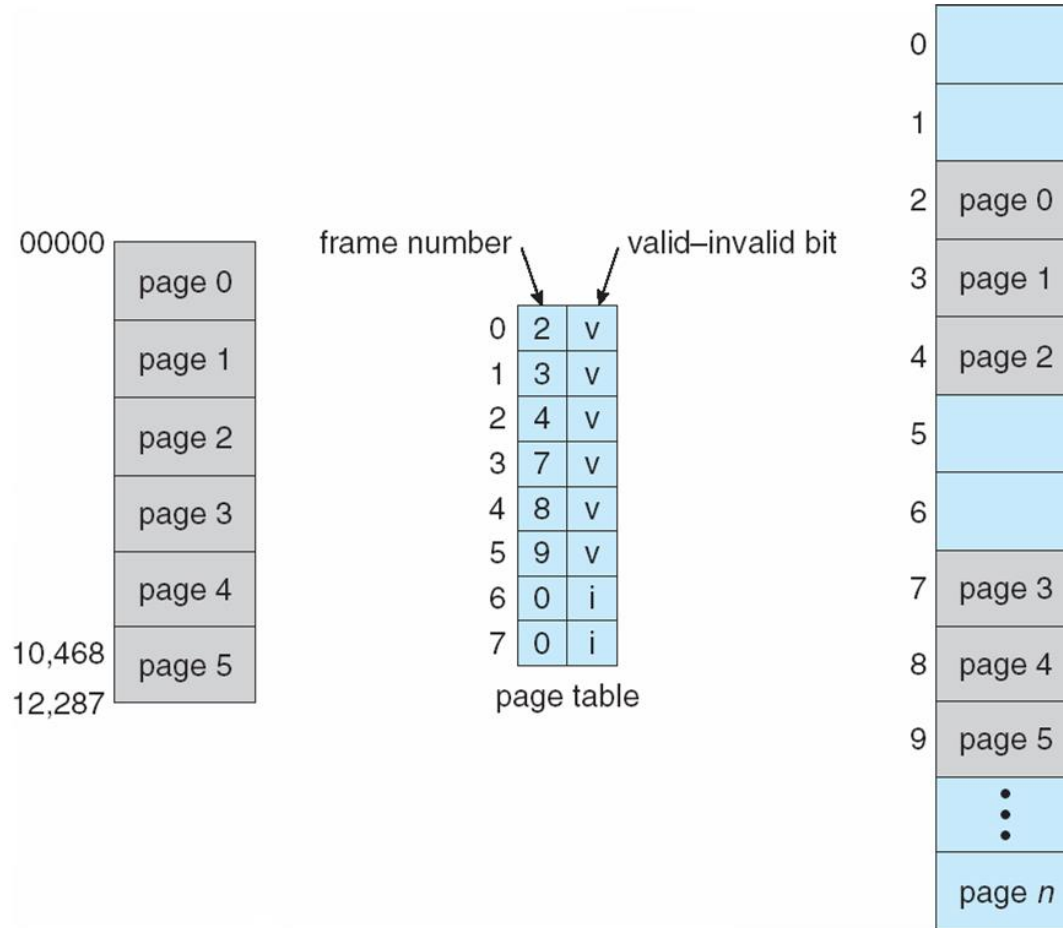
---

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
  - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
  - “invalid” indicates that the page is not in the process’ logical address space
  - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel





# Valid (v) or Invalid (i) Bit In A Page Table





# Shared Pages

---

## ■ Shared code

- One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed

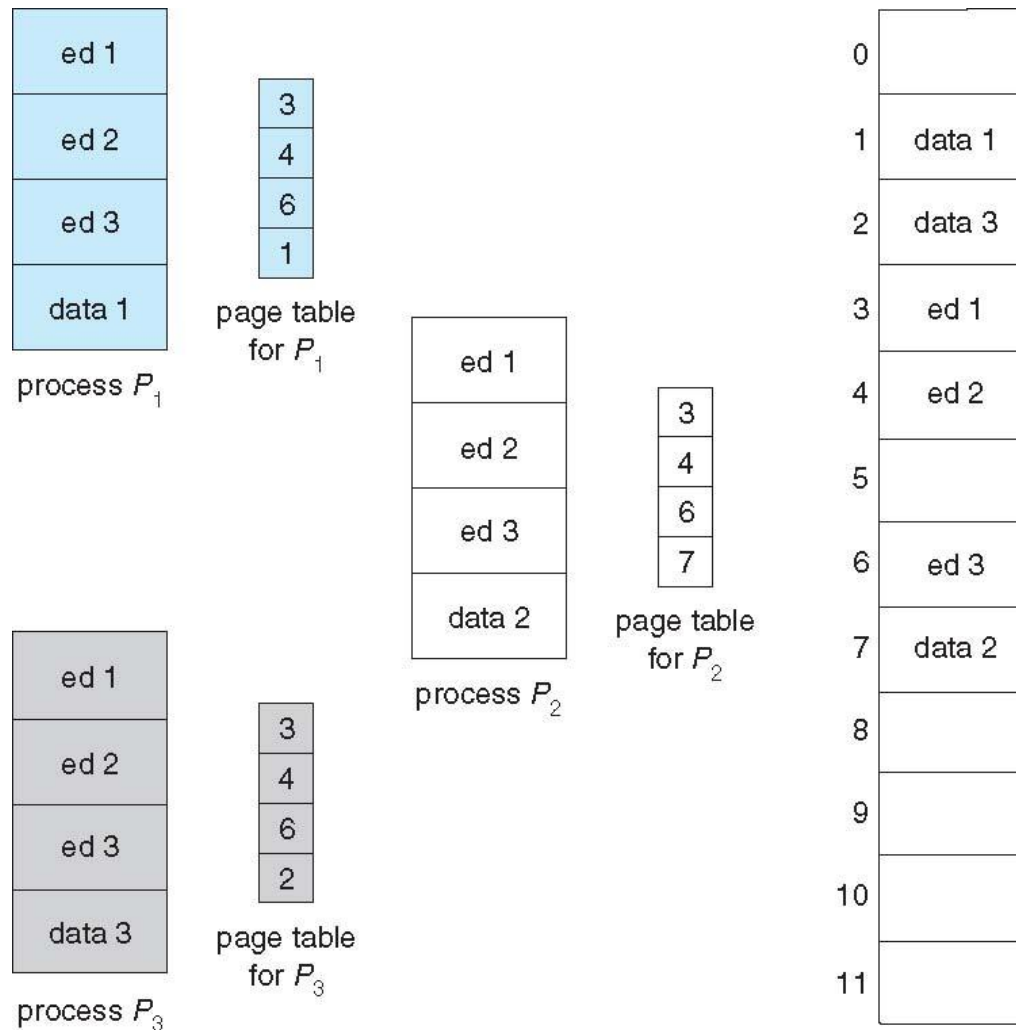
## ■ Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space





# Shared Pages Example





# Structure of the Page Table

- Memory structures for paging can get huge using straightforward methods
  - Consider a 32-bit logical address space as on modern computers
  - Page size of 4 KB ( $2^{12}$ ): offset is 12bits
  - Page table would have 1 million entries ( $2^{32} / 2^{12}$ )
  - If each entry is 4 bytes\*  $2^{20}$  -> 4 MB of physical address space / memory for page table alone
    - ▶ That amount of memory used to cost a lot
    - ▶ Don't want to allocate that contiguously in main memory
- Hierarchical Paging
  - Process not used 1G memory at a time
- Hashed Page Tables
  - Use only subset of space: grow need more p
- Inverted Page Tables
  - No need to allocate all pages:
  - make valid bit set only to few





# Hierarchical Page Tables

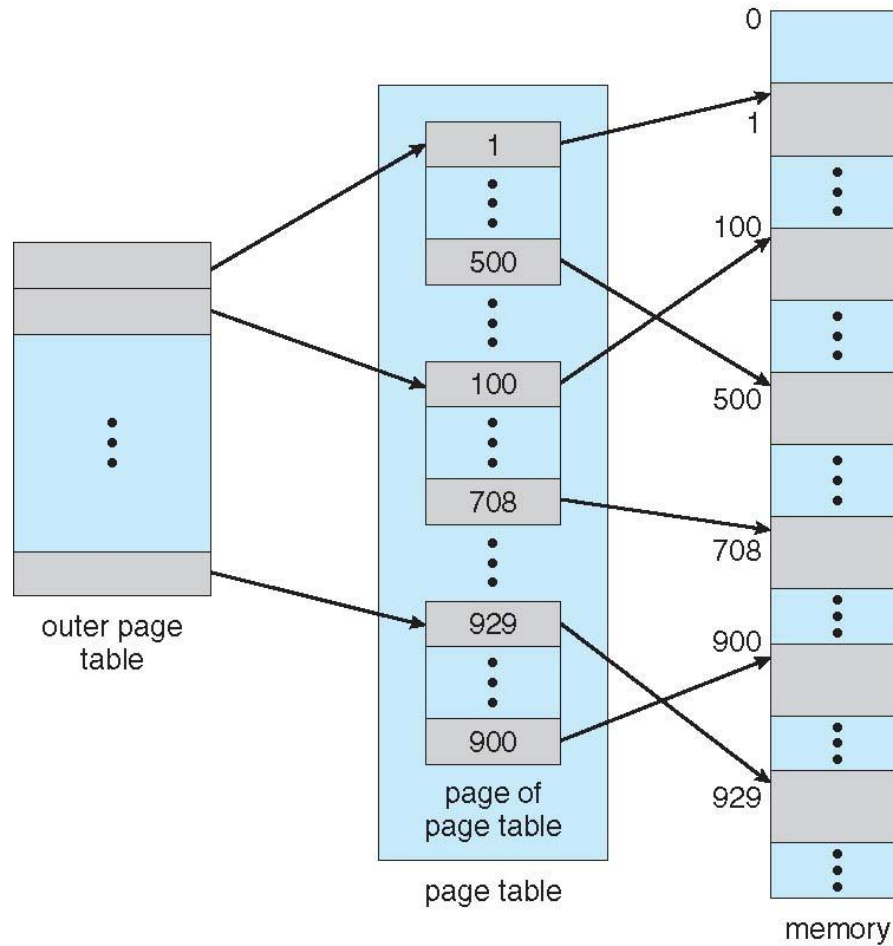
---

- Idea: Don't have one layer of page table
- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table
- Land line number example: 080 28603569





# Two-Level Page-Table Scheme





# Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
  - a page number consisting of 22 bits
  - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
  - a 12-bit page number
  - a 10-bit page offset
- Thus, a logical address is as follows:

page number		page offset
$p_1$	$p_2$	$d$
12	10	10

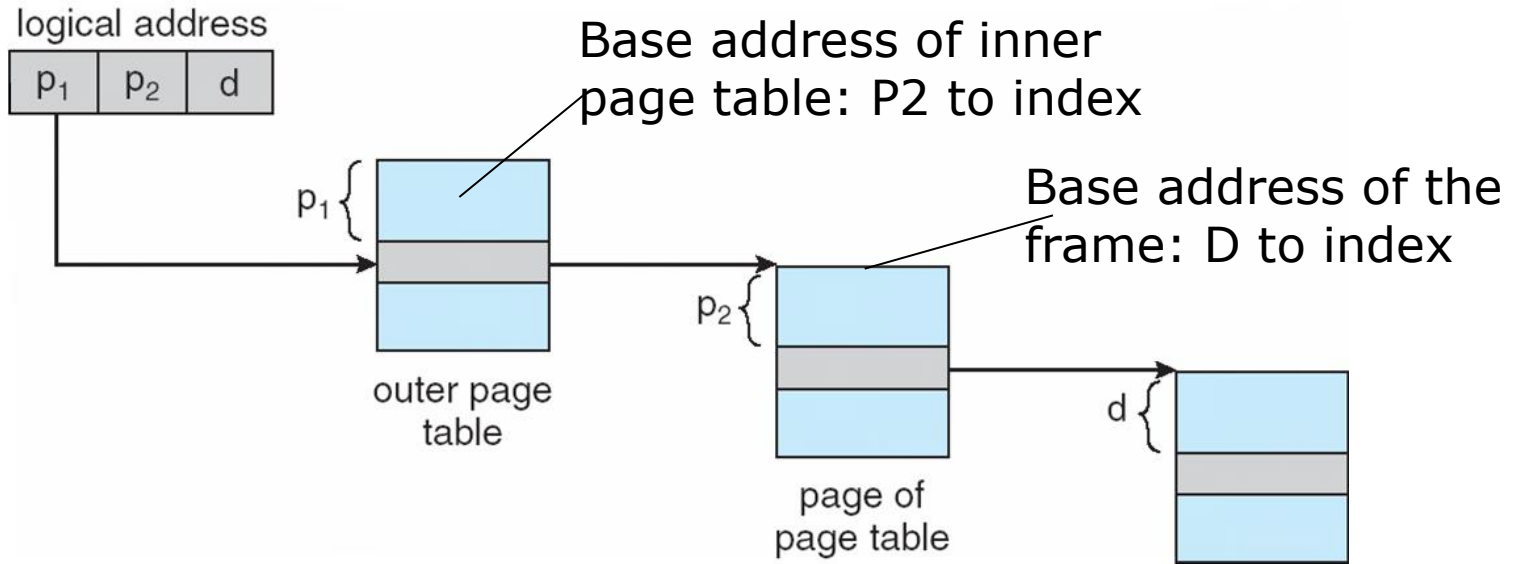
- where  $p_1$  is an index into the outer page table, and  $p_2$  to index into inner page table and offset to find actual content
- Known as **forward-mapped page table**








# Address-Translation Scheme





# 64-bit Logical Address Space

- Even two-level paging scheme not sufficient
- If page size is 4 KB ( $2^{12}$ )
  - Then page table has  $2^{52}$  entries
  - If two level scheme, inner page tables could be  $2^{10}$  4-byte entries
  - Address would look like



- Outer page table has  $2^{42}$  entries or  $2^{44}$  bytes
- One solution is to add a  $2^{\text{nd}}$  outer page table
- But in the following example the  $2^{\text{nd}}$  outer page table is still  $2^{34}$  bytes in size
  - ▶ And possibly 4 memory access to get to one physical memory location





# Three-level Paging Scheme

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12





# Hashed Page Tables

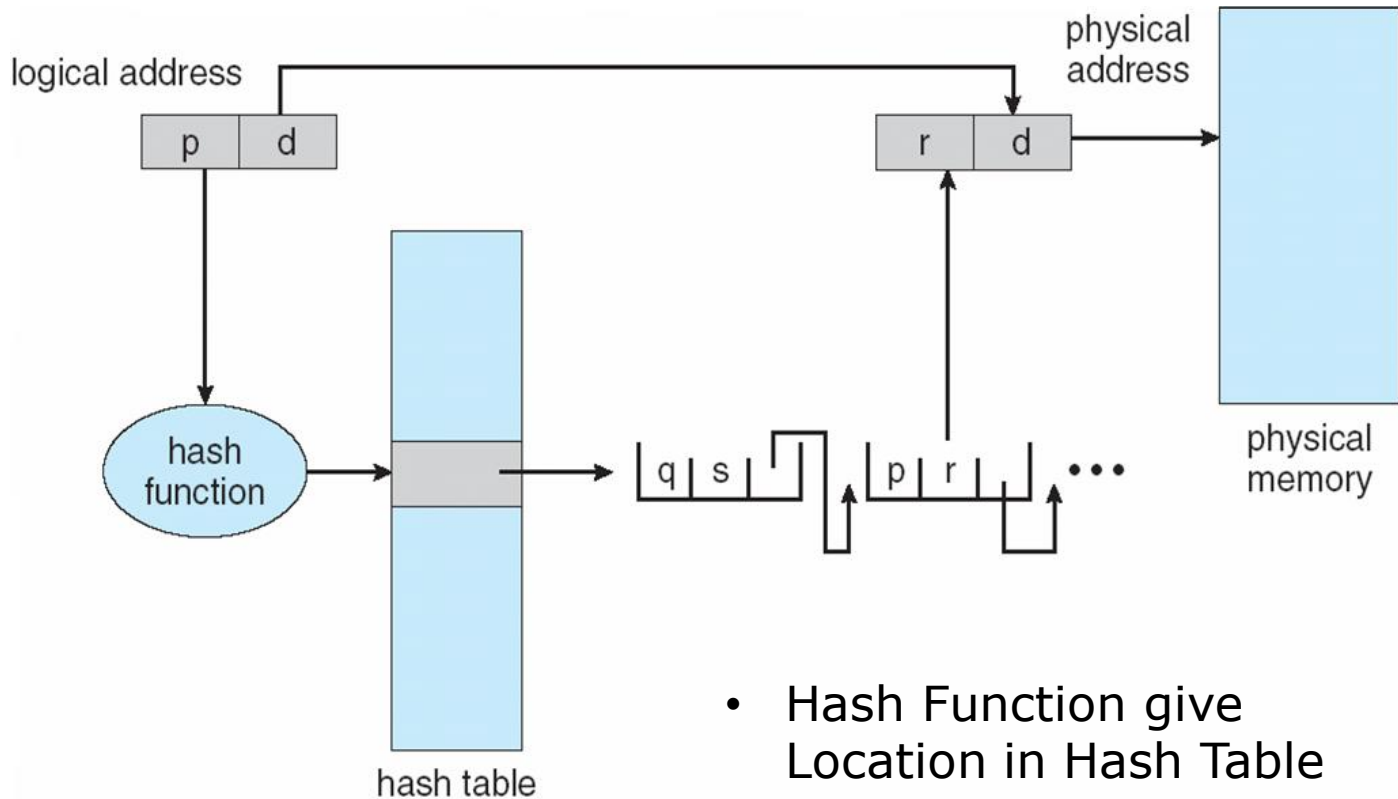
---

- Common in address spaces  $> 32$  bits
- The virtual page number is hashed into a page table
  - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
  - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is **clustered page tables**
  - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
  - Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)






# Hashed Page Table



- Hash Function give Location in Hash Table that has PT entry – Chained Entry
- Match entry with page number





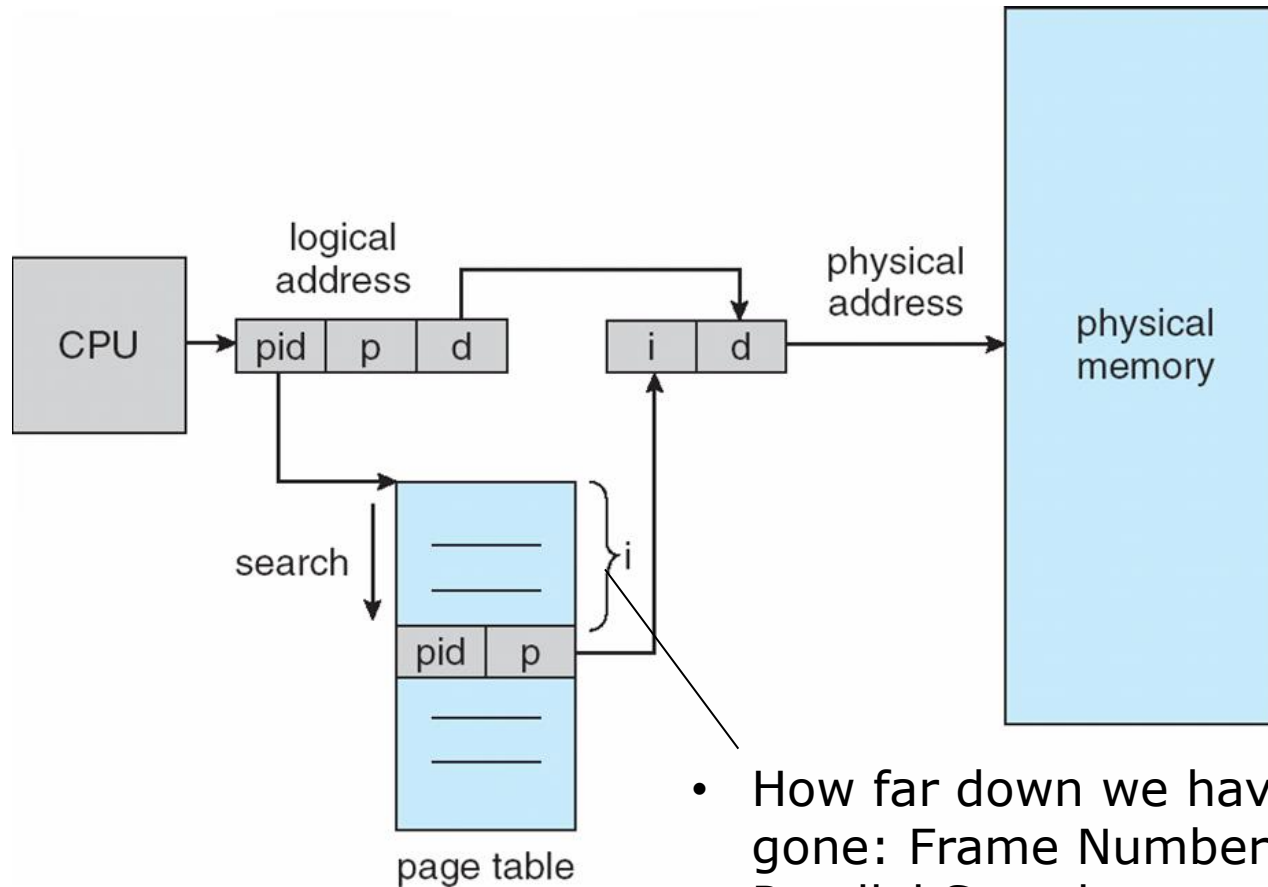
# Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries
  - TLB can accelerate access
- But how to implement shared memory?
  - One mapping of a virtual address to the shared physical address





# Inverted Page Table Architecture




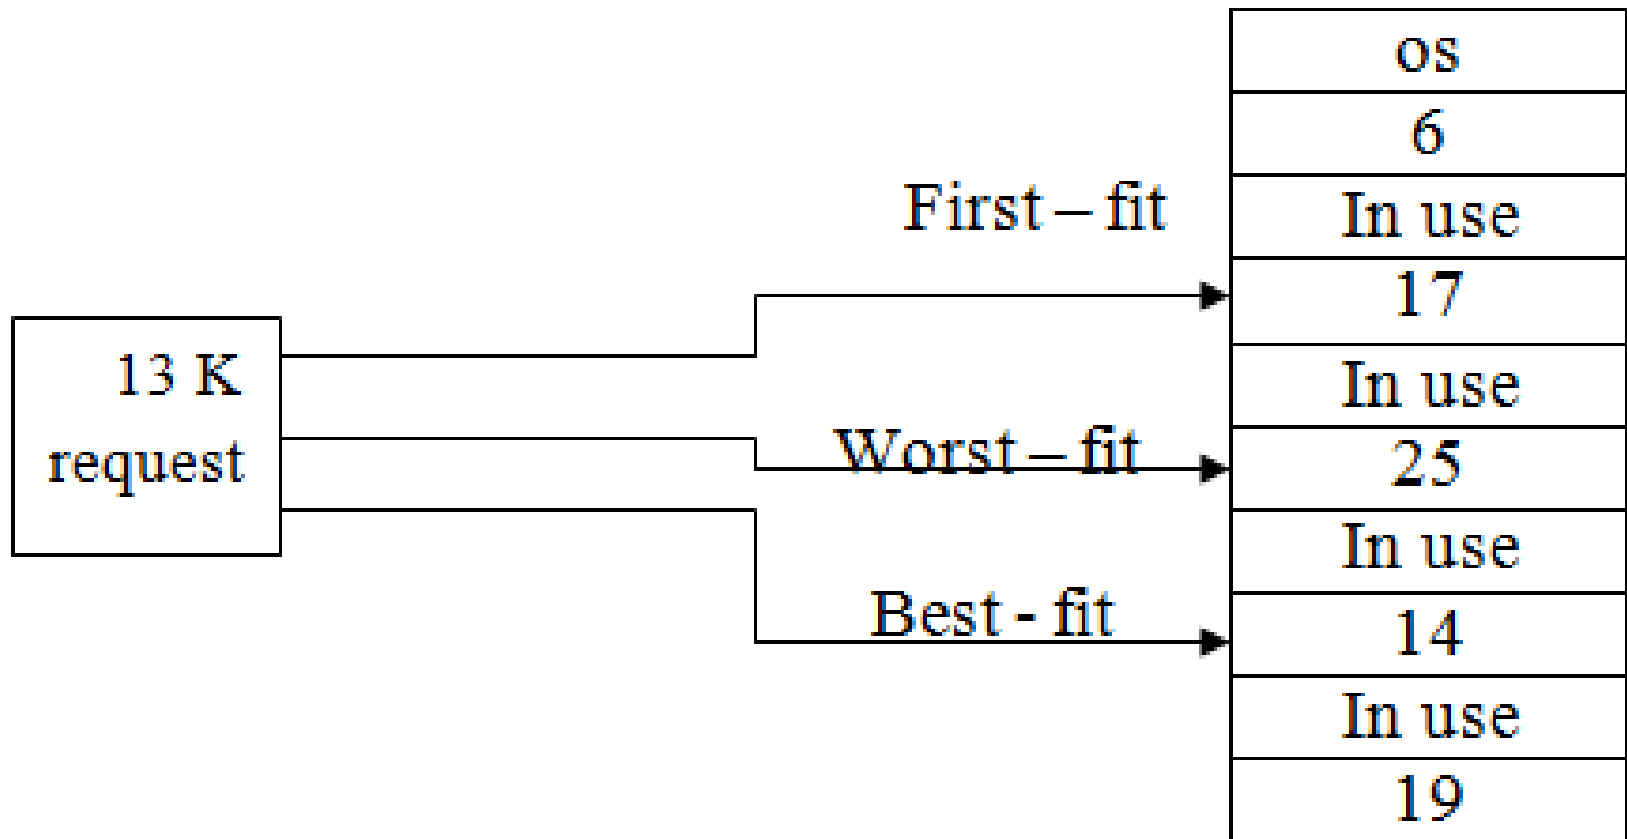
- How far down we have gone: Frame Number
- Parallel Search





# Numerical

Suppose that we have free segments with sizes: 6, 17, 25, 14, and 19. Place a program with size 13kB in the free segment using first-fit, best-fit and worst fit?







# Numerical

Consider a user program of logical address of size **6** pages and page size is **4** bytes.

The physical address contains **300** frames. The user program consists of **22** instructions **a, b, c, . . . u, v**.

Each instruction takes **1** byte. Assume at that time the free frames are **7, 26, 52, 20, 55, 6, 18, 21, 70, and 90**.

**Find the following?** (10 degrees)

- A) Draw the logical and physical maps and page tables?
- B) Allocate each page in the corresponding frame?
- C) Find the physical addresses for the instructions **m, d, v, r**?
- D) Calculate the fragmentation if exist?





# Numerical

## Answer Q6

0	. a . b . c . d
1	. e . f . g . h
2	. i . j . k . l
3	. m . n . o . p
4	. q . r . s . t
5	. u . v . . . .

Logical map

Page number	Frame number
0	7
1	26
2	52
3	20
4	55
5	6

page table

	Contents
6	. u . v . . . .
7	. a . b . c . d
20	. m . n . o . p
26	. e . f . g . h
52	. i . j . k . l
55	. q . r . s . t

Physical map

**The physical address = page size \* frame number + offset**

The physical address of m =  $4 * 20 + 0 = 80$

The physical address of d =  $4 * 7 + 3 = 31$

The physical address of v =  $4 * 6 + 1 = 25$

The physical address of r =  $4 * 55 + 1 = 221$

The external fragmentation = 0

The internal fragmentation = 2





# Numerical

Consider a program consists of five segments:  $S_0 = 600$ ,  $S_1 = 14 \text{ KB}$ ,  $S_2 = 100 \text{ KB}$ ,  $S_3 = 580 \text{ KB}$ , and  $S_4 = 96 \text{ KB}$ . Assume at that time, the available free space partitions of memory are 1200–1805, 50 – 150, 220-234, and 2500-3180. Find the following:


1. Draw logical to physical maps and segment table?
2. Allocate space for each segment in memory?
3. What are the addresses in physical memory for the following logical addresses: 0.580, (b) 1.17 (c) 2.66 (d) 3.82 (e) 4.20?






# Numerical

The logical map is:



Logical map





# Numerical

3. The physical addresses are

a) 0.580 → the physical address of 0.580 =  $1200 + 580 = 1780$ .

b) 1.17 → because  $d > \text{limit of } S1$ , the address is **wrong**.

c) 2.66 → the physical address of 2.66 =  $50 + 66 = 116$

d) 3.82 → the physical address is of 3.82 is =  $2500 + 82 = 2582$

e) 4.20 → the physical address 4.20 =  $3080 + 20 = 3100$





# Numerical

---

Consider a paging system with the page table stored in memory.

If a memory reference takes 200 nanoseconds, how long does a paged memory reference take?

If we add associative registers, and 75 percent of all page-table references are found in the associative registers, what is the effective memory reference time? (Assume that finding a page-table entry in the associative registers takes zero time, if the entry is there )





# Numerical

---

400 nanoseconds; 200 nanoseconds to access the page table and 200 nanoseconds to access the word in memory.

Effective access time =  $0.75 * (200 \text{ nanoseconds}) + 0.25 * (400 \text{ nanoseconds}) = 250 \text{ nanoseconds}$





# Numerical

Given memory partitions of 100K, 500K, 200K, 300K and 600K (in order), how would each of the First-fit, Best-fit and Worst-fit algorithms place processes of 212K, 417K, 112K and 426K (in order) ? Which algorithm makes the most efficient use of memory?

Answer:

First-fit

212K is put in 500K partition

417K is put in 600K partition

112K is put in 288K partition(new partition  $288K = 500K - 212K$ )

426K must wait

Best-fit

212K is put in 300K partition

417K is put in 500K partition

112K is put in 200K partition

426K is put in 600K partition

Worst-fit

212K is put in 600K partition

417K is put in 500K partition

112K is put in 388K partition

426K must wait







# Numerical

Consider a system has 4K pages of 512 bytes in size in the logical address space . Find the number of bits in logical address space.

Ans: Number of Pages: 4K

Page Size: 512B= $2^9$  Bytes

Number of Pages= Logical Address  
Space\*Page Size

=4K \*  $2^9$  Bytes

=  $2^{12}$  \*  $2^9$  Bytes

= $2^{21}$





# Numerical

Consider a system has 4K pages of 512 bytes in size in the logical address space . Find the number of bits in logical address space.

Ans: Number of Pages: 4K

Page Size: 512B= $2^9$  Bytes

Number of Pages= Logical Address  
Space\*Page Size

=4K \*  $2^9$  Bytes

=  $2^{12}$  \*  $2^9$  Bytes

= $2^{21}$





# Numerical

---

If a process occupies a logical address space of 39KB in a paging system with a page size of 4KB, determine the internal fragmentation in kilobytes

Ans: No. of Pages required: 20  
internal fragmentation in 1 kilobytes





# Numerical

Consider a system with byte-addressable memory, 32-bit logical addresses, 4 kilobyte page size and page table entries of 4 bytes each. The size of the page table in the system in megabytes is \_\_\_\_\_.

Ans:

$$\begin{aligned}\text{Number of entries in page table} &= 2^{32} / 4\text{Kbyte} \\ &= 2^{32} / 2^{12} = 2^{20}\end{aligned}$$

$$\begin{aligned}\text{Size of page table} &= (\text{No. page table entries}) * (\text{Size of an entry}) \\ &= 2^{20} * 4 \text{ bytes} = 2^{22} = 4 \text{ Megabytes}\end{aligned}$$



# End of Chapter 7

---

